PIWIK PRO

ClickHouse Operator


ClickHouse

Column oriented database for online analytical processing

Column oriented database for online analytical processing

- Column oriented database
  - Row - optimized for reading and writing rows efficiently
  - **Column** - optimized for reading and computing on columns efficiently
- Online analytical processing
  - OLTP - online **transaction** processing
  - **OLAP** - online **analytical** processing

Name	**City**	**Age**
Kamil	Wrocław	30
Marek	Opole	20

Name	Kamil	Marek
City	Wrocław	Opole
Age	30	20

ClickHouse

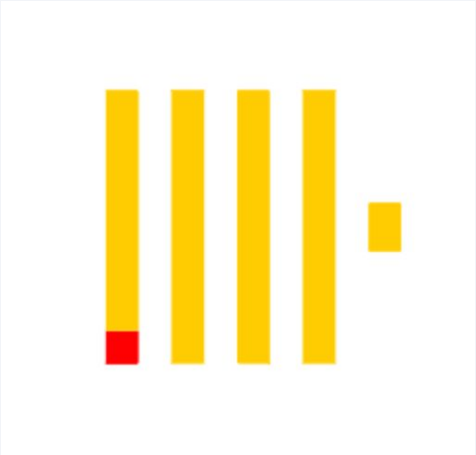

- Fast
- **Linear scalability** - scales to many petabytes
- Fault tolerant - **async master-master replication**
- Open source
- Developed by Yandex
- No transactions

Single server

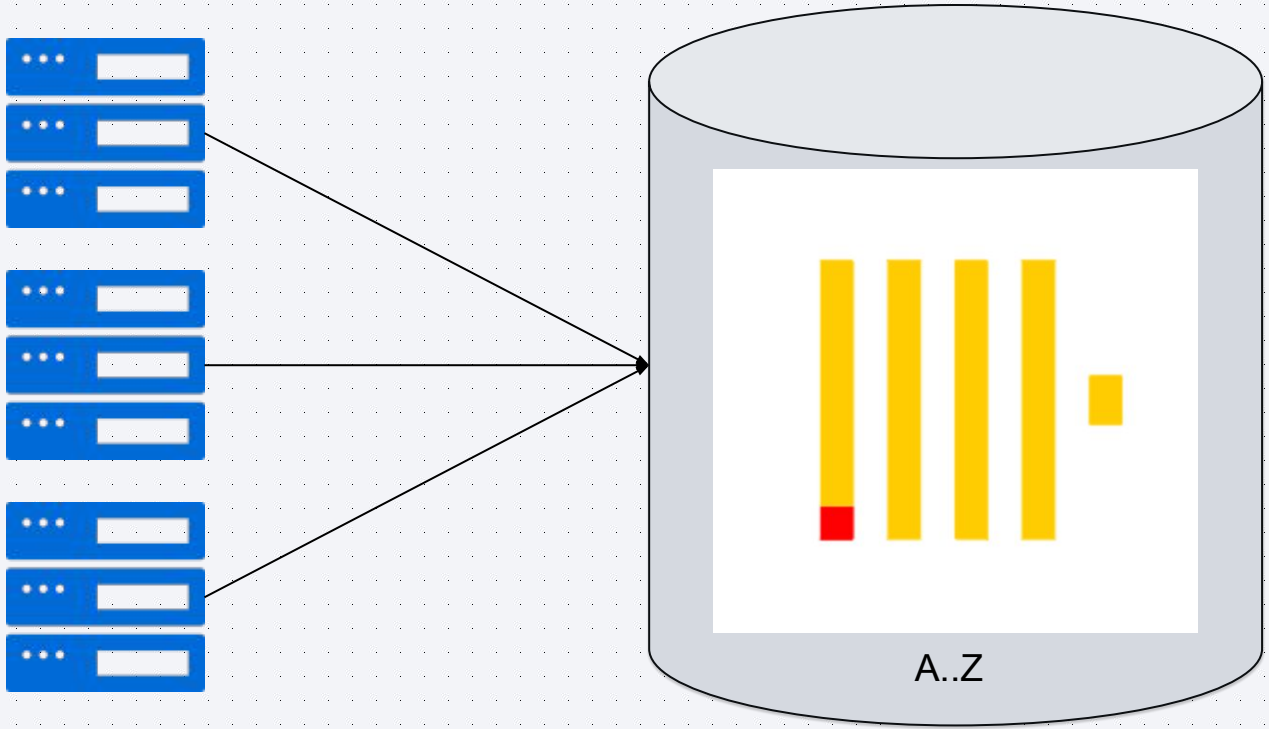

- Simple
- All applications use one server
- Not a real production setup

Problems:

- Limited scaling - only vertical
- Large dataset might not fit
- No High Availability

Sharding

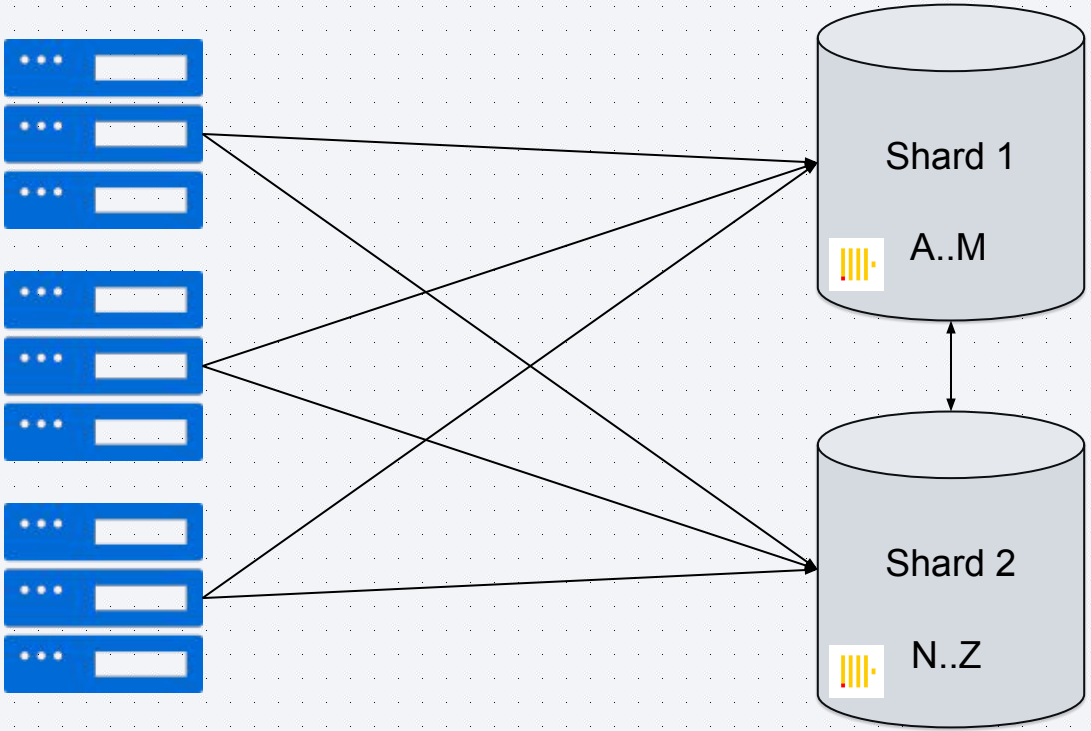

- Different part of data on each server
- Distributed tables
- Any sharding key

Problems:

- ✓ Horizontal and vertical scaling
- ✓ Large dataset might not fit
- No High Availability
- Load distribution on application side

Sharding

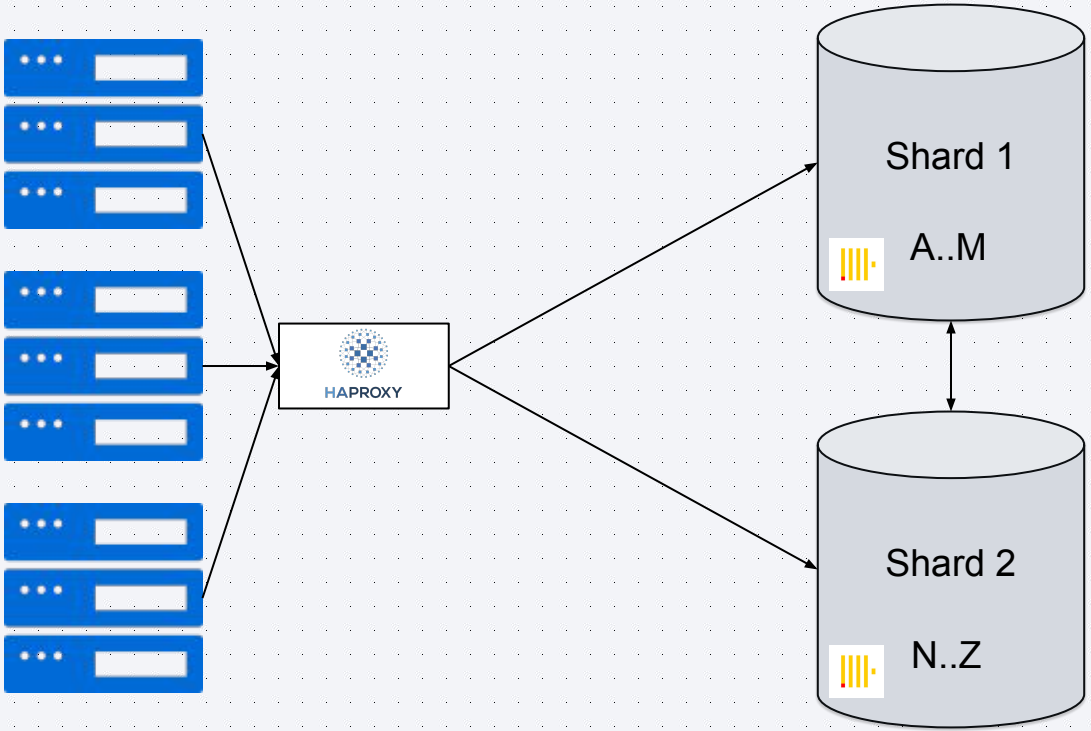

- Different part of data on each server
- Distributed tables
- Any sharding key

Problems:

- Horizontal and vertical scaling
- Large dataset might not fit
- No High Availability
- ✓ Load distribution on application side

Replication

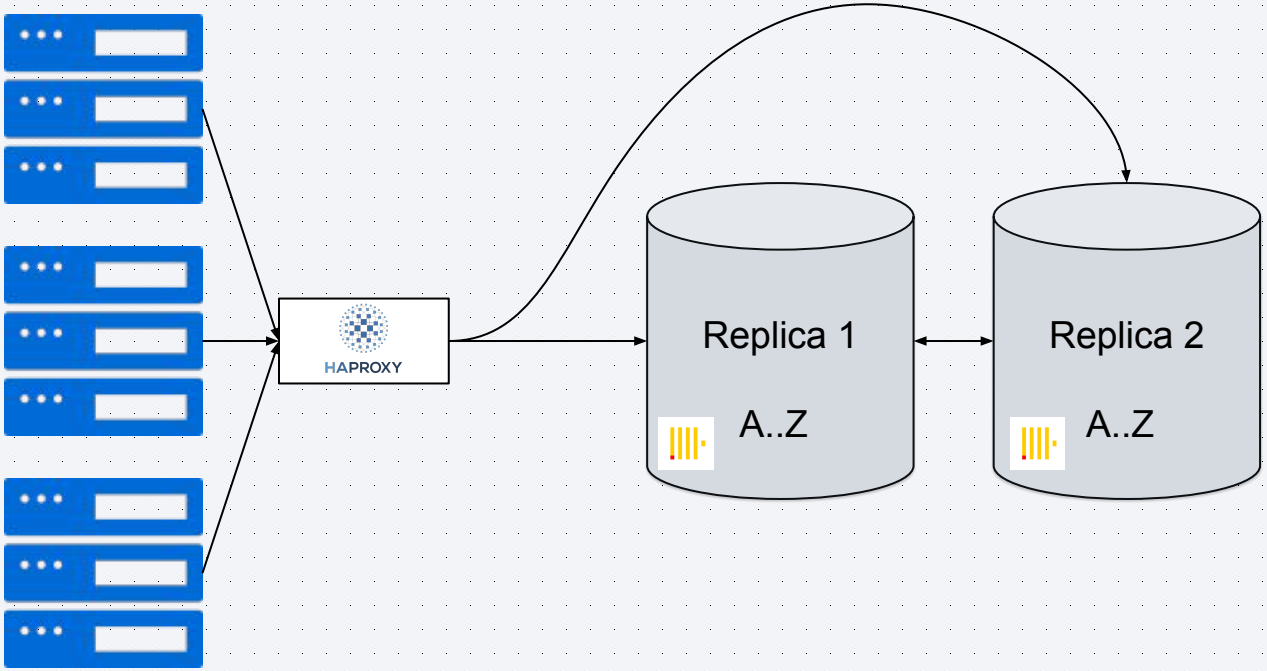

- Same data on all servers
- Each replica can handle reads
- **Master-master** replication - writes on all replicas

Problems:

✓ High Availability

Replication

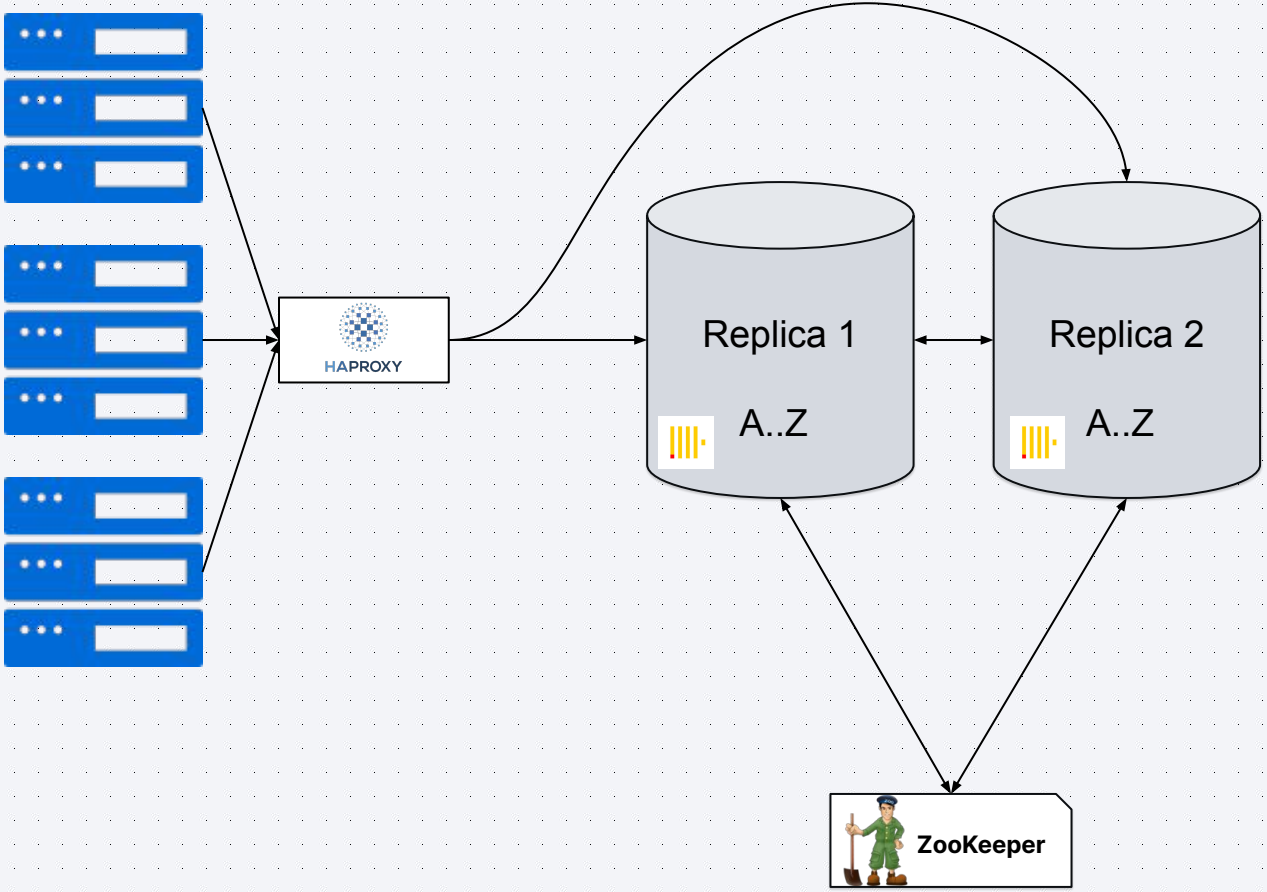

- Same data on all servers
- Each replica can handle reads
- **Master-master** replication - writes on all replicas
- ZooKeeper required for replication - only for inserts

Problems:

- High Availability

Replication in-depth - insert

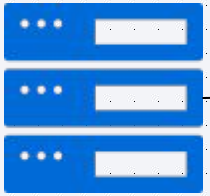

INSERT

Replica 1

Part: 202011_0_1_1

Part: 202011_2_2_0

Replica 2

Part: 202011_0_1_1

Part: 202011_2_2_0

Replication log:

Replication in-depth - insert

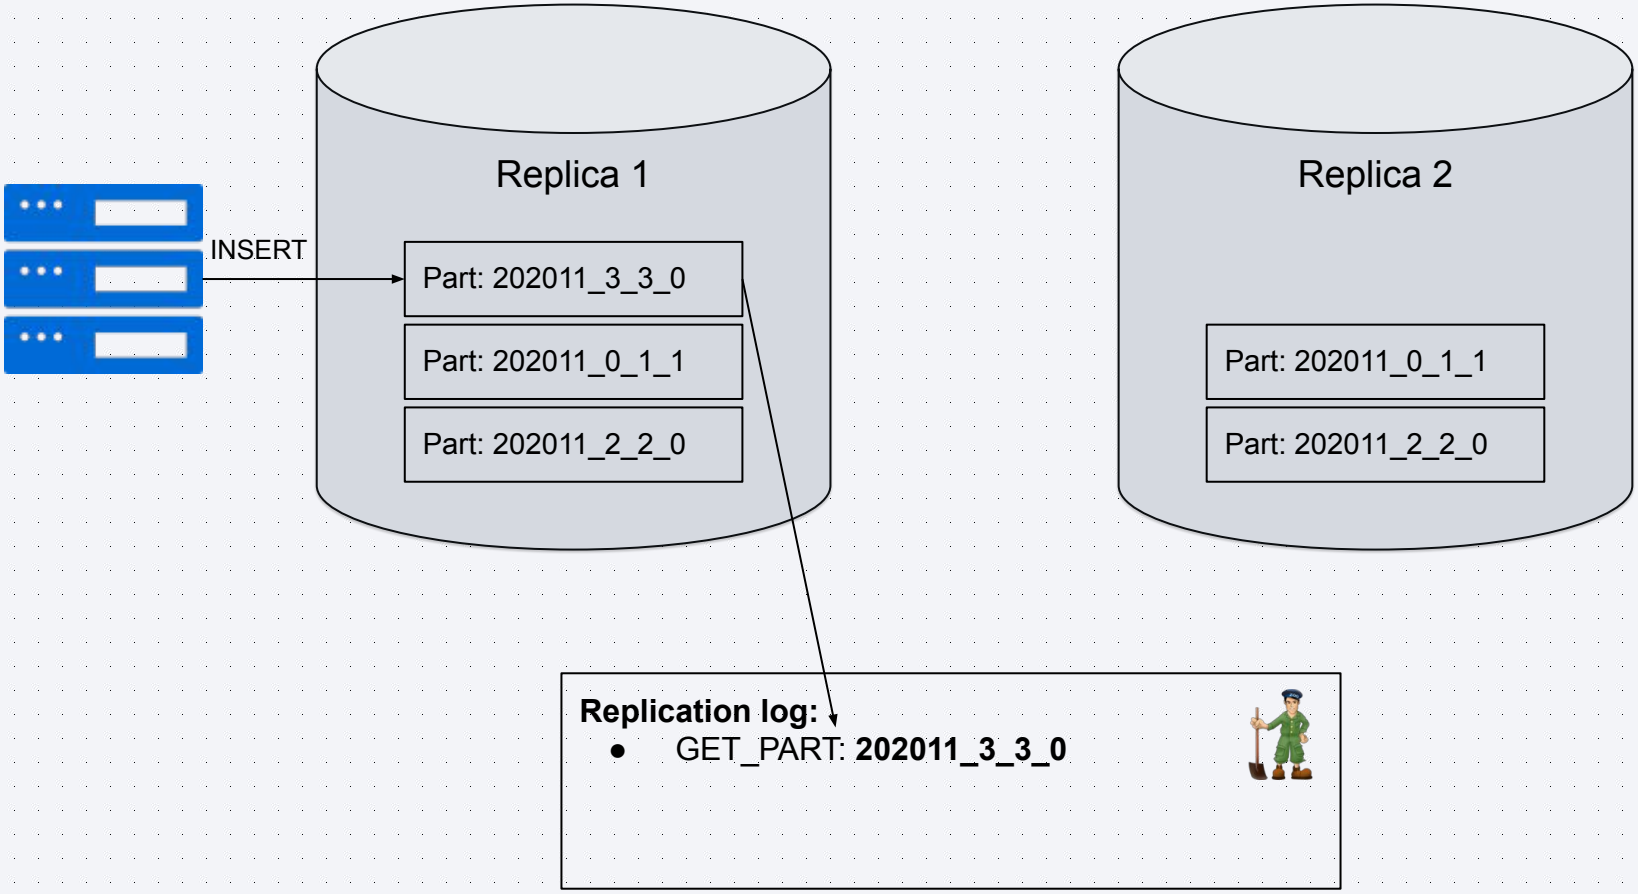

Replication in-depth - insert

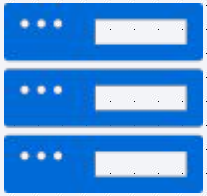

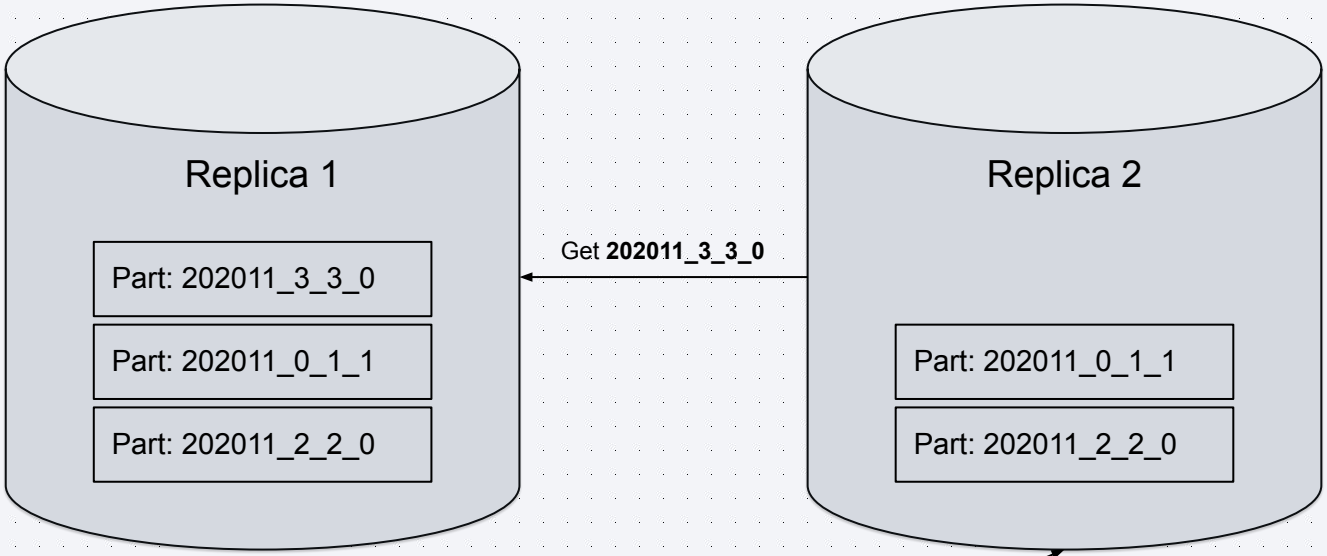

Replication log:

- GET_PART: **202011_3_3_0**

Replication in-depth - insert

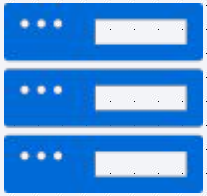

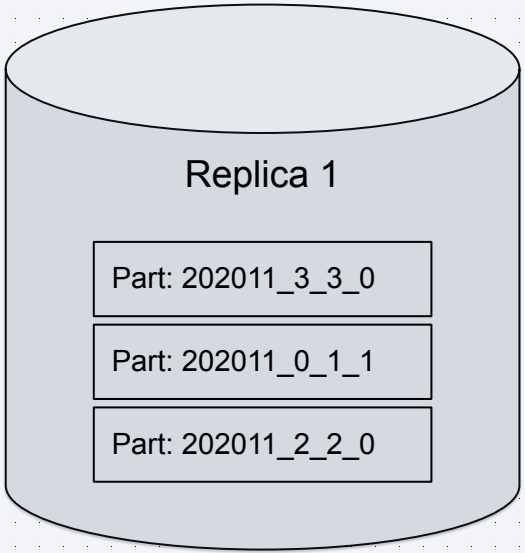

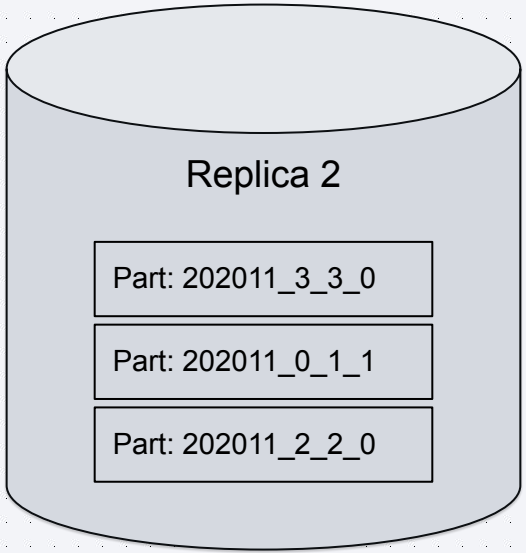

Replication log:

- GET_PART: **202011_3_3_0**

Replication in-depth - merging parts

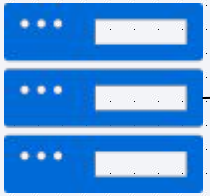

Replica 1

Part: 202011_3_3_0

Part: 202011_0_1_1

Part: 202011_2_2_0

Replica 2

Part: **202011_0_3_2**

Replication log:

- MERGE_PARTS: **202011_0_3_2**
  (202011_0_1_1, 202011_2_2_0, 202011_3_3_0)

Replication in-depth - merging parts

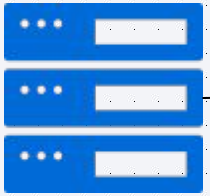

Replica 1

Part: **202011_0_3_2**

Replica 2

Part: 202011_0_3_2

Replication log:

- MERGE_PARTS: **202011_0_3_2**
  (202011_0_1_1, 202011_2_2_0, 202011_3_3_0)

Replication

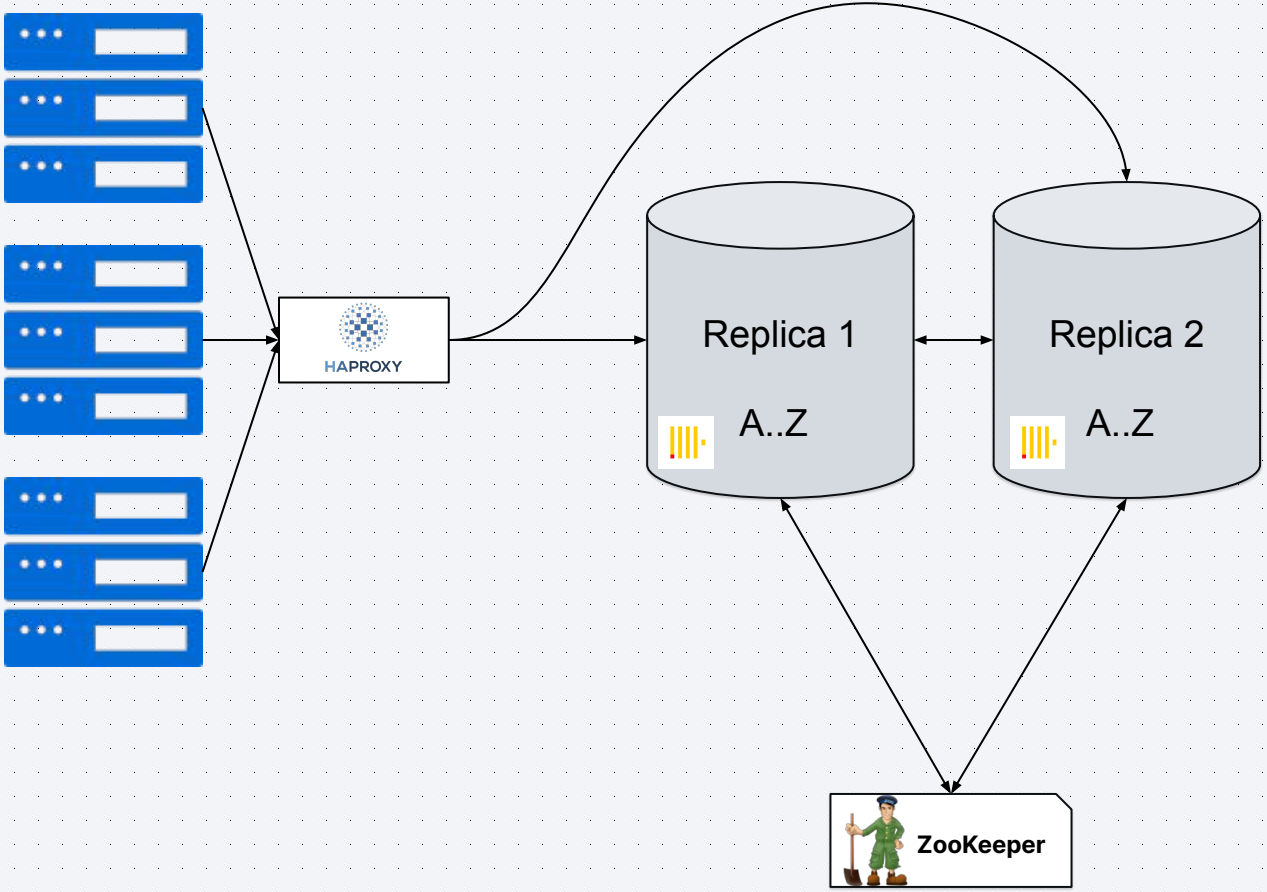

- Same data on all servers
- Each replica can handle reads
- Master-master replication - writes on all replicas
- ZooKeeper required for replication - only for inserts

Problems:

- High Availability
- High Availability for ZooKeeper

Replication

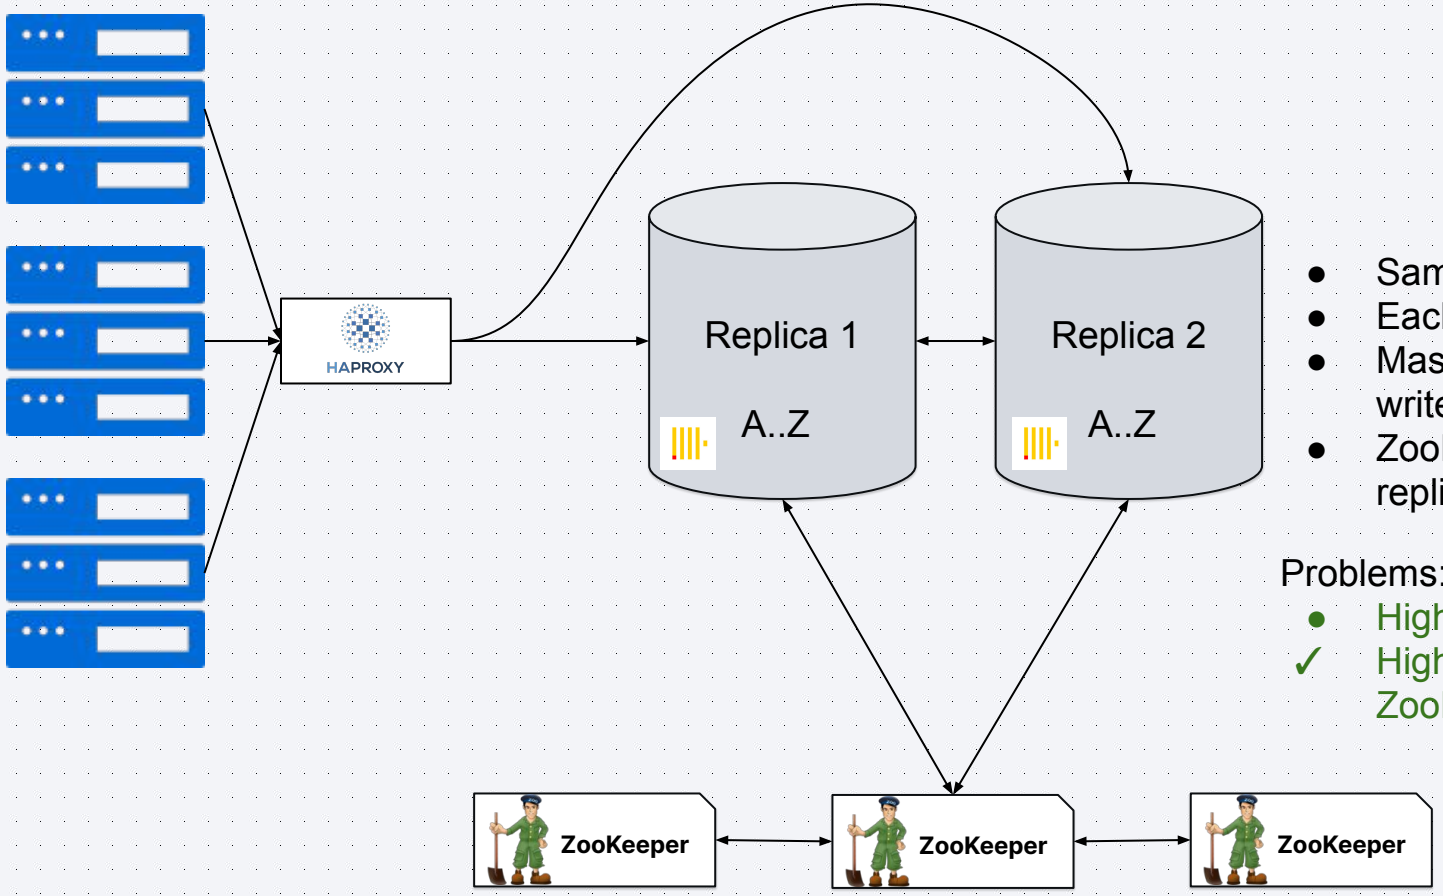

- Same data on all servers
- Each replica can handle reads
- Master-master replication - writes on all replicas
- ZooKeeper required for replication - only for inserts

Problems:

- High Availability
- ✓ High Availability also for ZooKeeper

Sharding & Replication

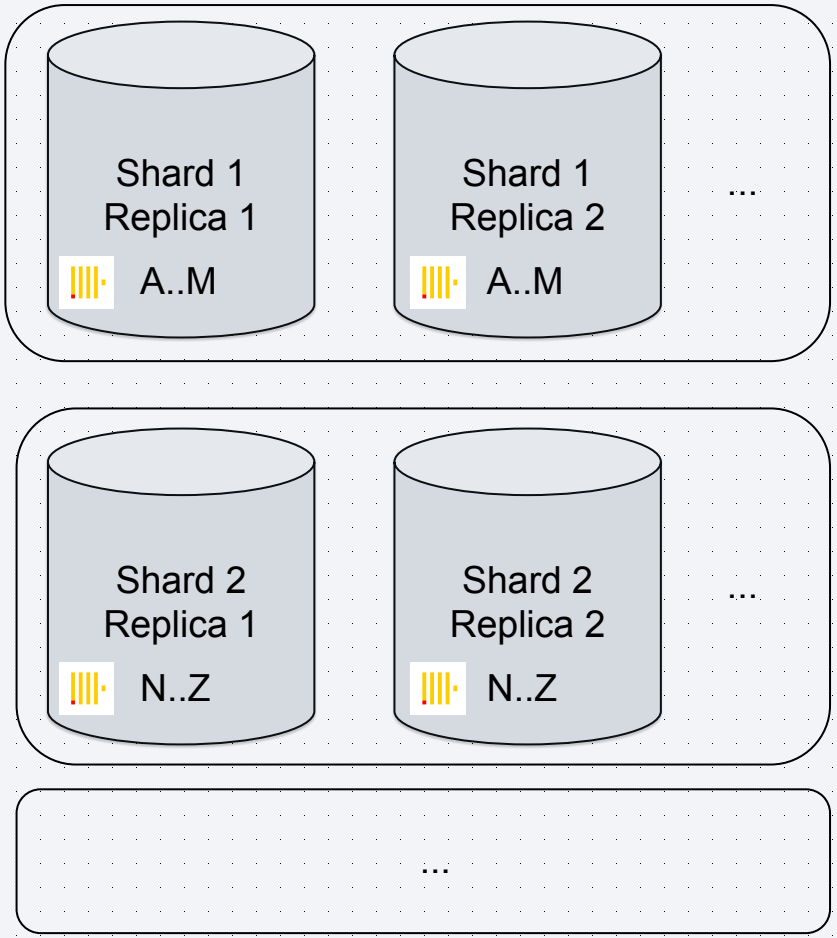

- Horizontal and vertical scaling
- High availability
- Master-master replication - writes and reads on all nodes
- Up to petabytes of data

Enterprise setup

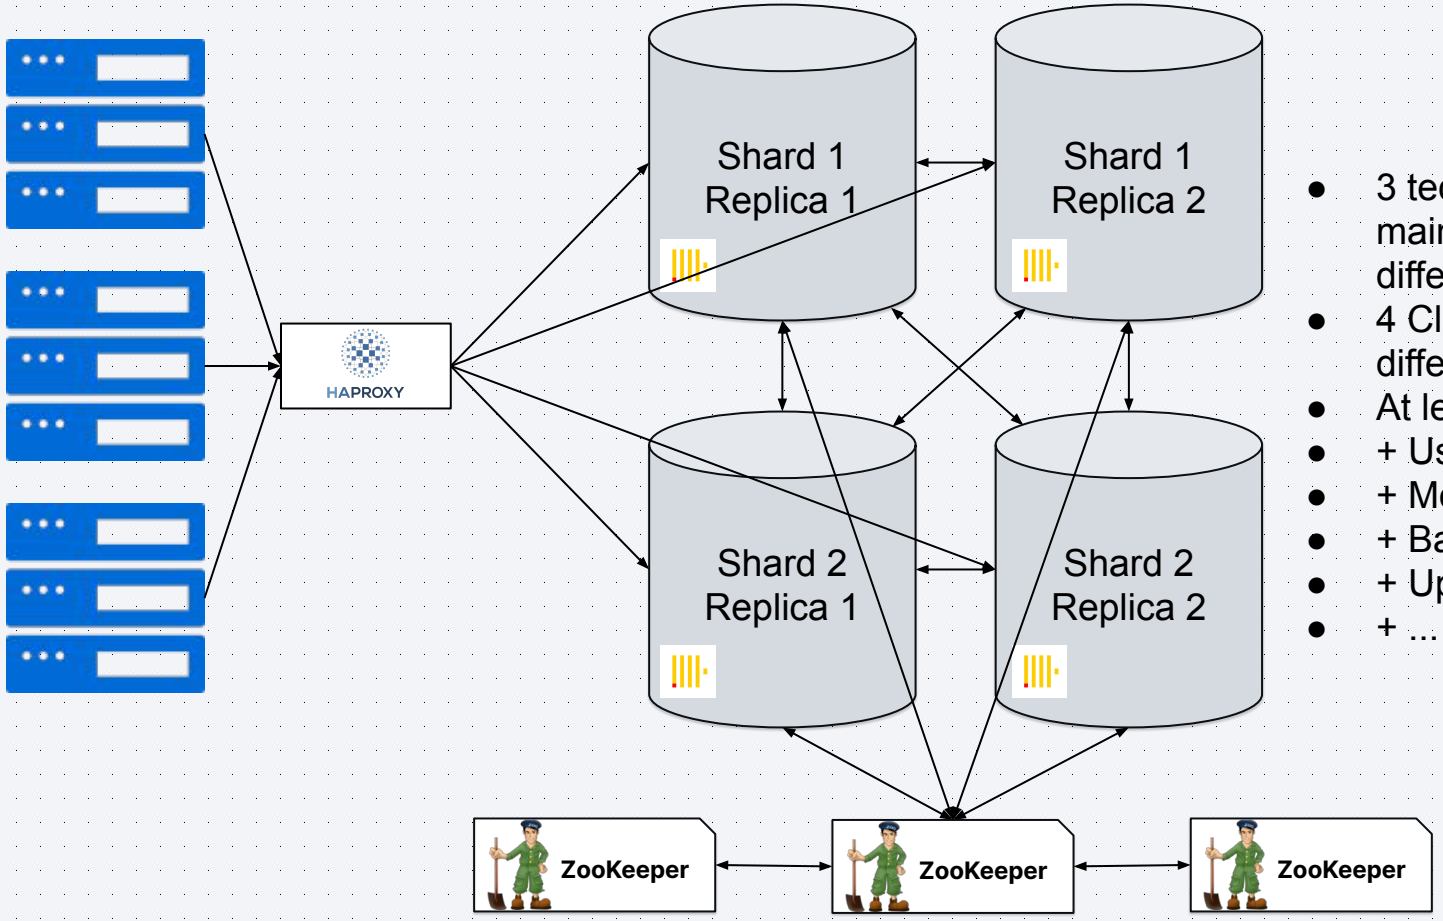

- 3 technologies - each maintained and scaled differently
- 4 ClickHouse servers in different roles
- At least 3 Zookeepers
- + Users mgmt
- + Monitoring
- + Backups
- + Upgrades
- + ...

Kubernetes (k8s)

How to run k8s locally

- minikube (minikube.sigs.k8s.io)
- kind (kind.sigs.k8s.io)
- k3d (k3d.io)

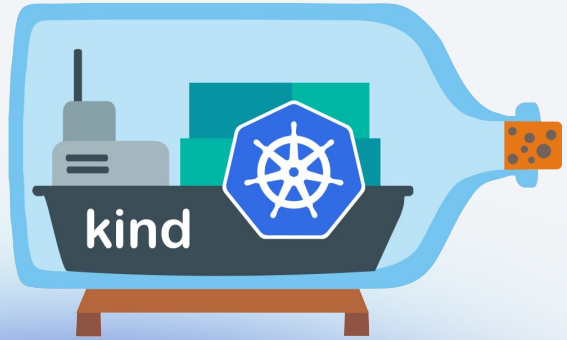

K8s operators

Operators in promil

- Creating MySQL database and users
- Creating Minio buckets
- Creating and configuration RabbitMQ queues
- Taking and restoring backups
- Deploying ClickHouse i Elasticsearch clusters
- Encryption of secrets (Sealed Secrets)

K8s operator - DIY

- Operator SDK (go, ansible, helm)
- Kopf (python)

Altinity ClickHouse Operator

Who is behind it?

- Altinity
  - Enterprise support for ClickHouse
  - Clickhouse feature development
- Founders:
  - Alexander Zaitsev
  - Peter Zaitsev and Vadim Tkachenko
- Great webinars on YouTube - check them out!

ClickHouseInstallation

Custom Resource

```yaml
apiVersion: "clickhouse.altinity.com/v1"
kind: "ClickHouseInstallation"
metadata:
  name: "test-05"
spec:
  defaults:
    templates:
      dataVolumeClaimTemplate: default
      podTemplate: clickhouse:20.7
   configuration:
    zookeeper:
      nodes:
      - host: zookeeper.zoo1ns
    clusters:
      - name: replicated
        layout:
          shardsCount: 2
          replicasCount: 2
  templates:
    volumeClaimTemplates:
      - name: default
        spec:
          resources:
            requests:
              storage: 500Mi
    podTemplates:
      - name: clickhouse:20.7
        spec:
          containers:
            - name: clickhouse-pod
              image: yandex/clickhouse-server:20.7
```

Database in container

- **Easier maintenance** - operators can use the same tools for application and database
- **Rollup upgrade and patching** - safe upgrades with easy way to rollback
- **Autoscaling** - adding additional replicas on hot hours
- **More alike dev, staging and production setups** - more likely to identify issues before production deploy
- **Similar setup on each platform** - Cloud/On-prem compatibility

Database in container

- Containers are meant to be **stateless**
- **Maintaining persistence** - after container shutdown data needs to be attached to another container (even on another node) in identical form
- **Availability** - Containers can be restarted at any time by oriestrator causing downtime
- Container will add an additional **performance overhead**
- **Security** - breach in other container may lead to compromising data
- **Maturity** - Is Kubernetes mature enough for running production databases?

Database in container

Containers are meant to be **stateless**
Data mounted in container

Maintaining persistence - after container shutdown data needs to be attached to another container (even on another node) in identical form
Assured with PersistentVolumes and Azure Managed Disks (on cloud)

Availability - Containers can be restarted at any time by oriestrator causing downtime
Production setups always with HA

Container will add an additional **performance overhead**
Yes, but how much?

Security - breach in other container may lead to compromising data
Separate k8s namespace and NetworkPolicies

Maturity - Is Kubernetes mature enough for running production databases?
Altinity sells ClickHouse on their k8s in SaaS model

Performance container vs on host

benchmarks

CPU usage

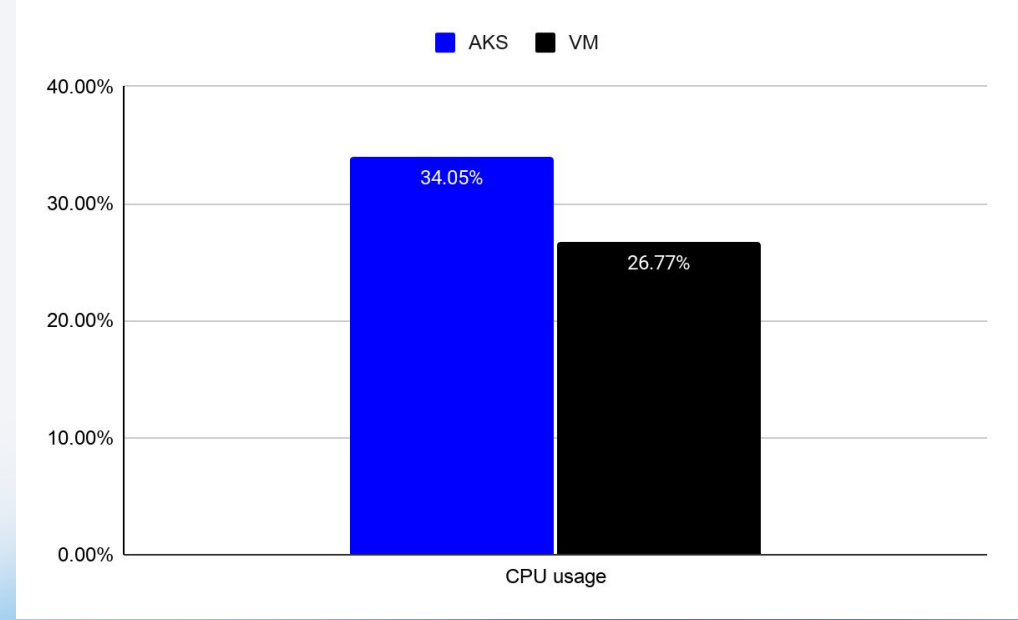

Monitoring

Prometheus

Grafana

- ClickHouse Operator + Prometheus & Grafana
- Metrics, Alerts and Dashboards provided by Clickhouse Operator
- Automatically adjust metrics collection based on cluster settings

Monitoring

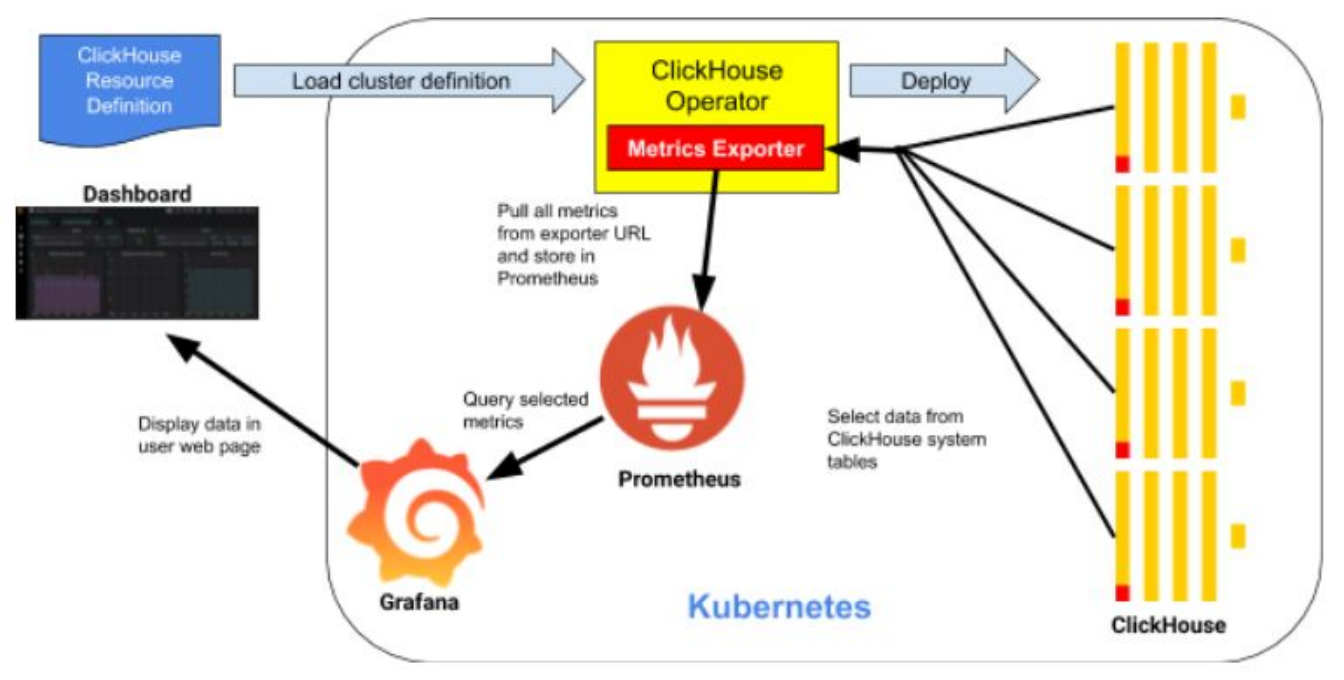

Source: https://altinity.com/wp-content/uploads/2020/07/Arch-Screenshotfrom2020-04-1517-15-27.png

Monitoring demo

Finito